

EXPERIMENT - 8

IMPLEMENTATION OF LEARNING ALGORITHMS FOR AN APPLICATION

AIM

To implement learning algorithms for an application

ALGORITHM

1. Linear regression is finding the best equation of line.
2. Independent and dependent variables are extracted
3. Split the data into training and test datasets. we import train and test from sklearn package
4. The next step is to fit the simple linear regression to training data set.
5. The next step is to predict the test results using the linear regression prediction package.
6. We then visualise the training set results by plotting the salary training set as graph. Salary is taken on the Y-axis and years of experience as the X-axis.

- 7 The next step is to find the residual accuracy
lower the value higher will be the accuracy
- 8 Calculate the errors including mean absolute error,
Mean squared error and root mean squared error.
- 9 Stop

FORMULA

Mean absolute error (MAE) - Distance between any data point
and best fit lines

Mean squared error (MSE) - Summation of square of all
distances between any data point

Root Mean Squared Error (RMSE) - Square of MSE

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$RMSE = \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{N}}$$

RESULT

Learning algorithm for salary dataset was implemented



linear_regression.ipynb ☆

File Edit View Insert Runtime Tools Help Last saved at 00:58

Comment Share ⚙



+ Code + Text

Connect Editing



```
[ ] import datetime

import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
%matplotlib inline

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score

[ ] dataset = pd.read_csv("Car_Dataset.csv")
dataset.head(5)
```

	Unnamed: 0	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type	Mileage	Engine	Power	Seats	New_Price
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	First	26.6 km/kg	998 CC	58.16 bhp	5.0	NaN
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	First	19.67 kmpl	1582 CC	126.2 bhp	5.0	NaN



linear_regression.ipynb ☆

File Edit View Insert Runtime Tools Help Last saved at 00:58

Comment Share ⚙



+ Code + Text

Connect Editing



```
[ ] dataset = pd.read_csv("Car_Dataset.csv")
dataset.head(5)
```

	Unnamed: 0	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type	Mileage	Engine	Power	Seats	New_Price
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	First	26.6 km/kg	998 CC	58.16 bhp	5.0	NaN
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	First	19.67 kmpl	1582 CC	126.2 bhp	5.0	NaN
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	First	18.2 kmpl	1199 CC	88.7 bhp	5.0	8.61 Lakh
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	First	20.77 kmpl	1248 CC	88.76 bhp	7.0	NaN
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Second	15.2 kmpl	1968 CC	140.8 bhp	5.0	NaN

```
[ ] X_train, X_test, y_train, y_test = train_test_split(dataset.iloc[:, :-1],
                                                        dataset.iloc[:, -1],
                                                        test_size = 0.3,
                                                        random_state = 42)
```

```
[ ] X_train.info()
```



linear_regression.ipynb ☆

File Edit View Insert Runtime Tools Help Last saved at 00:58

Comment Share



+ Code + Text

Connect ▾



```
[ ] X_train, X_test, y_train, y_test = train_test_split(dataset.iloc[:, :-1],
                                                    dataset.iloc[:, -1],
                                                    test_size = 0.3,
                                                    random_state = 42)
```

```
[ ] X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4213 entries, 4201 to 860
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   Unnamed: 0            4213 non-null  int64  
 1   Name                  4213 non-null  object  
 2   Location              4213 non-null  object  
 3   Year                  4213 non-null  int64  
 4   Kilometers_Driven    4213 non-null  int64  
 5   Fuel_Type             4213 non-null  object  
 6   Transmission         4213 non-null  object  
 7   Owner_Type           4213 non-null  object  
 8   Mileage               4212 non-null  object  
 9   Engine                4189 non-null  object  
10   Power                 4189 non-null  object  
11   Seats                 4185 non-null  float64 
12   New_Price             580 non-null   object  
dtypes: float64(1), int64(3), object(9)
memory usage: 460.8+ KB
```



linear_regression.ipynb ☆

File Edit View Insert Runtime Tools Help Last saved at 00:58

Comment Share ⚙



+ Code + Text

Connect ▾ Editing



```
[ ] X_train = X_train.iloc[:, 1:]
    X_test = X_test.iloc[:, 1:]
```

```
[ ] X_train["Name"].value_counts()
```

```
Mahindra XUV500 W8 2WD      35
Maruti Swift VDI             31
Maruti Ritz VDI             26
Hyundai i10 Sportz          25
Maruti Swift Dzire VDI       24
..
Skoda Laura L and K AT       1
Honda Amaze S Diesel         1
Nissan Micra XE               1
Renault KWID Climber 1.0 MT  1
Ford Endeavour 2.2 Titanium AT 4X2  1
Name: Name, Length: 1592, dtype: int64
```

```
[ ] make_train = X_train["Name"].str.split(" ", expand = True)
    make_test = X_test["Name"].str.split(" ", expand = True)
```

```
[ ] X_train["Manufacturer"] = make_train[0]
    X_test["Manufacturer"] = make_test[0]
```

linear_regression.ipynb

File Edit View Insert Runtime Tools Help Last saved at 00:58

+ Code + Text

Connect Editing

```
[ ] X_train.drop("Name", axis = 1, inplace = True)
X_test.drop("Name", axis = 1, inplace = True)

[ ] X_train.drop("Location", axis = 1, inplace = True)
X_test.drop("Location", axis = 1, inplace = True)

[ ] curr_time = datetime.datetime.now()
X_train["Year"] = X_train["Year"].apply(lambda x : curr_time.year - x)
X_test["Year"] = X_test["Year"].apply(lambda x : curr_time.year - x)

[ ] X_train["Kilometers_Driven"]

4201    77000
4383    19947
1779    70963
4020    115195
3248    58752
...
3772    27000
5191     9000
5226   140000
5390    76414
860     98000
Name: Kilometers_Driven, Length: 4213, dtype: int64
```

linear_regression.ipynb

File Edit View Insert Runtime Tools Help Last saved at 00:58

+ Code + Text

Connect Editing

```
[ ] mileage_train = X_train["Mileage"].str.split(" ", expand = True)
mileage_test = X_test["Mileage"].str.split(" ", expand = True)

X_train["Mileage"] = pd.to_numeric(mileage_train[0], errors = 'coerce')
X_test["Mileage"] = pd.to_numeric(mileage_test[0], errors = 'coerce')

[ ] print(sum(X_train["Mileage"].isnull()))
print(sum(X_test["Mileage"].isnull()))

1
1

[ ] X_train["Mileage"].fillna(X_train["Mileage"].astype("float64").mean(), inplace = True)
X_test["Mileage"].fillna(X_train["Mileage"].astype("float64").mean(), inplace = True)

[ ] cc_train = X_train["Engine"].str.split(" ", expand = True)
cc_test = X_test["Engine"].str.split(" ", expand = True)
X_train["Engine"] = pd.to_numeric(cc_train[0], errors = 'coerce')
X_test["Engine"] = pd.to_numeric(cc_test[0], errors = 'coerce')

bhp_train = X_train["Power"].str.split(" ", expand = True)
bhp_test = X_test["Power"].str.split(" ", expand = True)
X_train["Power"] = pd.to_numeric(bhp_train[0], errors = 'coerce')
X_test["Power"] = pd.to_numeric(bhp_test[0], errors = 'coerce')
```

linear_regression.ipynb

File Edit View Insert Runtime Tools Help Last saved at 00:58

Comment Share

+ Code + Text

Connect Editing

```
[ ] X_train["Engine"].fillna(X_train["Engine"].astype("float64").mean(), inplace = True)
    X_test["Engine"].fillna(X_train["Engine"].astype("float64").mean(), inplace = True)

X_train["Power"].fillna(X_train["Power"].astype("float64").mean(), inplace = True)
X_test["Power"].fillna(X_train["Power"].astype("float64").mean(), inplace = True)

X_train["Seats"].fillna(X_train["Seats"].astype("float64").mean(), inplace = True)
X_test["Seats"].fillna(X_train["Seats"].astype("float64").mean(), inplace = True)

[ ] X_train.drop(["New_Price"], axis = 1, inplace = True)
    X_test.drop(["New_Price"], axis = 1, inplace = True)

[ ] X_train = pd.get_dummies(X_train,
                           columns = ["Manufacturer", "Fuel_Type", "Transmission", "Owner_Type"],
                           drop_first = True)

[ ] X_test = pd.get_dummies(X_test,
                           columns = ["Manufacturer", "Fuel_Type", "Transmission", "Owner_Type"],
                           drop_first = True)

missing_cols = set(X_train.columns) - set(X_test.columns)
for col in missing_cols:
    X_test[col] = 0
```

linear_regression.ipynb

File Edit View Insert Runtime Tools Help Last saved at 00:58

Comment Share

+ Code + Text

Connect Editing

```
[ ] missing_cols = set(X_train.columns) - set(X_test.columns)
    for col in missing_cols:
        X_test[col] = 0
    X_test = X_test[X_train.columns]

[ ] standardScaler = StandardScaler()
    standardScaler.fit(X_train)
    X_train = standardScaler.transform(X_train)
    X_test = standardScaler.transform(X_test)

[ ] linearRegression = LinearRegression()
    linearRegression.fit(X_train, y_train)
    y_pred = linearRegression.predict(X_test)
    r2_score(y_test, y_pred)

0.7008908549416727

[ ] rf = RandomForestRegressor(n_estimators = 100)
    rf.fit(X_train, y_train)
    y_pred = rf.predict(X_test)
    r2_score(y_test, y_pred)

0.8882005989619215
```